

Introduction to Computational and Algorithmic Thinking

NATURAL LANGUAGE PROCESSING AND REGULAR EXPRESSIONS



Announcements

This lecture: Natural Language Processing and Regular Expressions

Reading: Read Chapter 10 of Conery

Acknowledgement: Some of this lecture slides are based on CSE 101 lecture notes by Prof. Kevin McDonald at SBU

The Turing Test

- Alan Turing (1912-1954) was an English mathematician who laid some of the important theoretical groundwork of computer science
- In addition to other topics, Turing was interested in the idea of computers being able to think as human beings do
- He devised what he called the **imitation game**, now known as the **Turing test**
- A human judge (the interrogator) engages in an online chat with another person and a computer, but isn't told which is which

The Turing Test

- The interrogator asks the human being and the computer a series of questions to carry on a conversation
- If the interrogator cannot tell which is the person and which is the computer, then the computer has passed the Turing Test because it is simulating human intelligence
- So the Turing Test touches on two important areas of computer science: **artificial intelligence** and **natural language processing** (NLP)
- [Google AI passes Turing test](#)

ELIZA: The First (?) Chatbot

- ELIZA is a program first developed in the 1960s at MIT that attempts to converse in English with a human participant
 - More specifically, the program “pretends” to be a psychotherapist, so its responses often consist of questions that use words the user entered
- Based on an internal set of rules that analyze what the user types, the software attempts to simulate a natural conversation
- For example, one kind of rule will look for key words in the human’s entry (H) and respond with generic sentences (C). Here’s an example:
 - **H: My father wouldn't buy me a puppy.!**
 - **C: Tell me more about your family.!**

ELIZA: The First (?) Chatbot

- Since ELIZA is simulating a psychotherapist, its rules pay special attention to words like "dream" and various words that connote emotions:
 - **H: I'm sad I never got a puppy.!**
 - **C: Can you explain what made you sad you never got a puppy?!**
- These rules have limits. Consider the following exchange, which the software could generate:
 - **H: I'm sad we never had a chance to play fetch.!**
 - **C: I am sorry to hear you are sad we never had a chance to play fetch.!**

ELIZA: The First (?) Chatbot

- The software's intelligence is limited by its rules, or **script**
- The program should become better at conversing as more rules are added to its script
- The simplest rule would simply be that the program repeats or *echoes* the words entered by the user:

```
from PythonLabs.ElizaLab import Eliza  
Eliza.run()
```

- Sample run:

H: hello!

C: hello!

H: are you imitating me?!

C: are you imitating me?!

bye!

ELIZA: The First (?) Chatbot

- To make things more interesting, let's load the DOCTOR script, which simulates a psychotherapist.
- We could create our own scripts to simulate other kinds of conversations.

How do you do. Please tell me your problem.

H: I am afraid of spiders.

C: Is it because you are afraid of spiders that you came to me?

H: Yes, it is.

C: You seem quite positive.

H: I am very happy.

C: How have I helped you to be happy?

H: Just by chatting with me.

C: I am not sure I understand you fully.

H: You are not very smart.

C: What makes you think I am not very smart?

H: bye

Sentence Patterns

- To find key words in a sentence we need to search through strings
- The **find** method for strings will get the job done
- Given a string to search through and a target substring, the **find** method will return the index of the first character of the substring
- If the substring is not found, the method returns **-1**
- If the substring appears in several places, the **find** method
- returns the index of the first instance
- Example:
s1 = 'I was afraid of the cow.'
s1.find('cow') # returns 20

Sentence Patterns

- Now we will see that the method will return **-1** if we search for a substring that is not in the string:
- Example:
s1 = 'I am afraid of the dark.'
s1.find('cow') # returns -1!
- Now let's see what happens if we search for a substring that appears multiple times in a string:
s1 = 'I saw birds on the house, birds at school, birds everywhere!'
s1.find('birds') # returns 6
- See [eliza_functions.py](#) for these examples and other function calls on the next 20+ slides

Sentence Patterns

- The **find** method takes an optional argument that indicates at what index the search should begin
- For example, the first instance of “birds” in the string below is at index 6.
- If we start the search at index 7, the **find** method will find the next instance of “birds” in the string, which starts at index 26:

```
s1 = 'I saw birds on the house, birds at school, birds everywhere!'
```

```
s1.find('birds', 7) # returns 26
```

Sentence Patterns

- One issue with **find** is that does not have any notion of *words*
- For example, the letters of “cow” appear in the word “scowl”, among others
 - If we were trying to build a chatbot to have conversations about farm animals, it might incorrectly see the word “scowl” and think that the person had typed “cow”
- So to extract key words from a user’s entry, we actually need to look for *words*, not simply substrings
- To help address this problem we will use **regular expressions**, also called **regexes**, for short

Regular Expressions

- Regular expressions give us a formal ways of expressing **patterns** of characters we want to find in an input string
- The simplest kind of regular expression looks for a particular substring in an input string
- Initially we will not look at the details of how to write regular expressions, but rather use some capabilities from the ElizaLab
- The **Pattern** class in ElizaLab lets us create regular expressions using a user-friendly notation
- The **Pattern** object below can be used to detect when a sentence contains the word “cow”:

```
from PythonLabs.ElizaLab import Pattern  
p = Pattern('cow')
```

Regular Expressions

- Given a **Pattern** object, we can then provide possible responses the chatbot might produce when it sees a sentence that matches the pattern:

```
p.add_response('Tell me more about your farm')  
p.add_response('Go on')
```

- We can also give a list of responses when we create the **Pattern** object, instead of doing it later:

```
p = Pattern('lamb', ['I love lambs', 'How cute'])
```

Regular Expressions

- To see if a sentence matches the pattern, we call the **apply** method of the **Pattern** object
 - If the match is successful, a response string is returned
 - If not, the method returns **None**
- An example:
p.apply('I milked the cow and fed the chickens.')
 - Return value:
'Tell me more about your farm'

Regular Expressions

- The **apply** method is “smart” enough to look for words, not simply substrings
- For the example below, “scowled” will not trigger a match for the word “cow”, so **None** is returned:

p.apply('I scowled at the horse yesterday.')

Regular Expressions

- Regular expressions let us define patterns that will match several different words
 - We separate the words by vertical bars to form a **group**
 - This means we don't have to make a separate pattern for each word
- Below is a pattern that could be applied to any sentence containing “cow”, “pig” or “horse”:
p = Pattern('cow | pig | horse', ['Really?', 'Go on'])
p.apply('The horse jumped the fence.')
- For this input, **apply** returns **'Really?'**

Decomposing Sentences into Parts

- Finding individual words in a sentence is useful, but it's not enough to understand the context of the word
- We need some way of decomposing a sentence into parts
- First, we will need to somehow save the part of the sentence that matches a group in pattern so it can be used later to make the response
- Each group in a pattern has **an associated variable name**, with **\$1** identifying the first group, **\$2** for the second group, and so on
- We can then use these variable names in responses
- This is easiest to understand by example (next slide)

Decomposing Sentences into Parts

```
p = Pattern('cow | pig | horse')  
p.add_response('You had a $1?')  
p.add_response('How many $1s were on the farm?')
```

- Now, call the **apply** method to try the new pattern:

```
p.apply('The horse jumped the fence.')
```

- Return value: '**You had a horse?**'

Decomposing Sentences into Parts

- We can start to build up pretty complex patterns that contain multiple groups, as in this example:

```
p = Pattern('I (like | love | adore) my  
            (cat | dog | ducks)')  
p.add_response('Why do you $1 your $2?')  
p.add_response('What about your $2 do you $1?')
```

- Call the **apply** method to try the new pattern:
- **p.apply('I adore my cat.')**
- Return value: **'Why do you adore your cat?'**

Decomposing Sentences into Parts

- This approach looks pretty good, but we would have to list every word we are interested in
- What if we wanted to write a pattern to anticipate everything a person might be afraid of?
 - We really couldn't list everything
- We could write a pattern containing one or more wildcards that match any piece of text
- A wildcard is written using a period followed by an asterisk

Decomposing Sentences into Parts

- Example of a wildcard expression:

```
p = Pattern('I am afraid of .*')  
p.add_response('Why are you afraid of $1?')  
p.apply('I am afraid of little green men')
```

- Return value:

```
'Why are you afraid of little green men?'
```

- This looks pretty good, but it can lead to some nonsensical results, as we'll see in the next example

Response Pre-/Post-processing

- The patterns we just wrote use a basic copy-and-paste strategy to transfer a string fragment from the input to the response, but can lead to silly responses:

```
p = Pattern('I am (.*)', ['Are you really $1?'])  
p.apply('I am happy to see you')
```

- Return value:

```
'Are you really happy to see you?'
```

- What we need to do after matching the input is to do some **postprocessing** to replace some words or phrases
- For example, the response above really should have been “Are you really happy to see *me*?”

Response Pre-/Post-processing

- Let's create a dictionary that will provide substitutions for pronouns that should be made for a pattern:

```
pronouns = { 'I': 'you', 'your': 'my' }
```

- Now let's try a few examples:

```
p.apply('I am sorry I dropped your  
computer', post=pronouns)
```

- Return value:

```
'Are you really sorry you dropped my computer?'
```

- Another example:

```
p.apply("I'm happy I lost", post=pronouns)
```

- Return value: **None** (uh-oh! What happened?)

Response Pre-/Post-processing

- The **apply** function can match **"I"**, but not **"I'm"**
- Let's define a dictionary that can recognize **contractions**
- We will use it to perform **preprocessing** on the input string before we attempt to make a match:
contractions = { "I'm": "I am" }
p.apply("I'm happy I lost",
pre=contractions, post=pronouns)
- Return value:
'Are you really happy you lost?'

Response Pre-/Post-processing

- When ElizaLab is used to load a script, the preprocessing and postprocessing dictionaries are available via **Eliza.pre** and **Eliza.post**, respectively:
- **Eliza.pre** has the value:

```
{"can't": 'can not', "don't": 'do not', "i'm": 'I am', "won't": 'will not', "you're": 'you are'}
```
- **Eliza.post** has the value:

```
{'am': 'are', 'are': 'am', 'i': 'you', 'me': 'you', 'my': 'your', 'myself': 'yourself', 'you': 'I', 'your': 'my', 'yourself': 'myself'}
```

Algorithm for Having a Conversation

- The summary so far:
 - A **Pattern** object is defined by a string that has a key word, groups of words, or wildcards
 - We can specify any number of response strings, and the object's **apply** method will return one of these strings if a sentence matches its pattern
 - Pattern-matching variables (**\$1**, **\$2**, etc.) along with preprocessing and postprocessing help the **apply** method extract parts of the input sentence and reuse them as part of the response
- Now, we need to take this functionality and develop an algorithm for carrying on a conversation

Algorithm for Having a Conversation

- One simple approach is to try patterns until we find one that matches an input sentence, but this means that patterns will always be applied in the same order
- To solve this problem we can assign a **priority** to each word
 - The creator of ELIZA anticipated that people would ask questions like “Are you a computer?” so he gave higher priority to words like “computer” and “machine” than others
- Another problem is that the same word (like “**I**”) can appear in multiple patterns - which one should be applied?
 - In ElizaLab, this problem is addressed by creating a **Rule** object, which is a collection of patterns that pertain to a particular word

Algorithm for Having a Conversation

- The **rule_for** method tells us which rules (if any) pertain to a particular word:

Eliza.rule_for('remember')

- Return value:

PythonLabs.ElizaLab.Rule [5]

'I remember (.*)'

'do you remember (.*)'

- The **5** indicates that this rule has high priority (higher number à higher priority)
- We can use a priority queue (from our study of Huffman coding) to keep track of the highest-priority rules

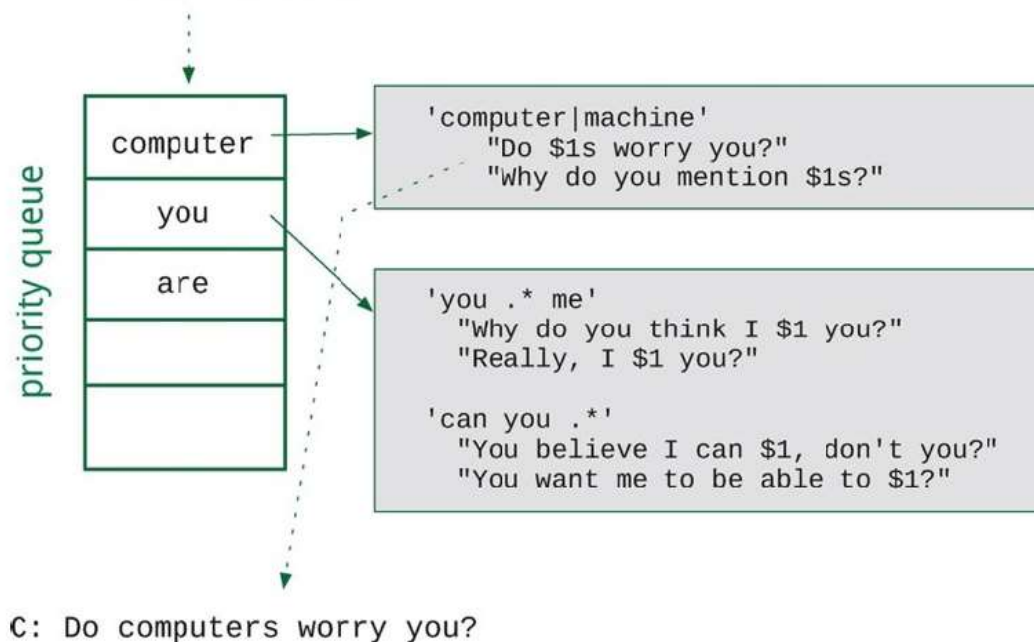
Algorithm for Having a Conversation

- A complete algorithm for transforming an input sentence into a response has three steps:
 1. initialization
 2. tokenization (borrowed from the SpamLab) and insertion of rules into priority queue
 3. rule processing

Algorithm for Having a Conversation

- The priority queue and rules relevant to this sentence are depicted in the figure below

H: *Are you a computer?*



Algorithm for Having a Conversation

```
from PythonLabs.Tools import tokenize
from PythonLabs.ElizaLab import *

def transform(sentence):
    queue = PriorityQueue()
    sentence = Eliza.preprocess(sentence, Eliza.pre)

    for word in tokenize(sentence):
        rule = Eliza.rule_for(word)
        if rule is not None:
            queue.insert(rule)

    while len(queue) > 0:
        rule = queue.pop()
        response = rule.apply(sentence, post=Eliza.post)
        if response is not None:
            return response

    return None
```

See [eliza_transform.py](#)

Algorithm for Having a Conversation

- Let's call the **transform** function a few times:

transform('Are you a computer?')

- Sample return values:

Do computers worry you?

Why do you mention computers?

**What do you think computers have to do
with your problem?**

**Don't you think computers can help
people?**

What about computers worries you?

ELIZA and the Turing Test

- Can a chatbot like ELIZA pass the Turing Test?
- One challenge in developing natural language processing systems is algorithmically assigning meanings to words in a sentence. This is the problem of **semantics**.
- Consider the following exchange:
H: That cow was huge.
C: How big was it?
- This kind of response is not possible with a rule-based algorithm like Eliza's because the software doesn't understand that "huge" and "big" have similar definitions.
- So we could give the program a dictionary and/or thesaurus

ELIZA and the Turing Test

- Giving the program a dictionary of word definitions is not enough, however. Consider the following exchange:

H: The cow fell into the canal.

C: Was it able to get out?

- Having the definitions of all the words in the human being's sentence will not empower a computer to **draw conclusions, make inferences, and apply other kinds of “real-world knowledge” or “common sense”**
- So, although natural language processing systems have been deployed successfully in systems like airline reservation systems and other interactive telephone-based systems, there is still a long way to go before a computer could consistently pass the Turing Test

Regular Expressions

“Real” Regular Expressions

- As we have seen, part of string processing in real applications (e.g., online forms in websites) often involves ensuring that a string matches a particular format, called a **pattern**
- We have been expressing patterns using the ElizaLab module, but these aren’t “real” regular expressions (although they are pretty close)
- Now, we will look at how **regular expressions** are actually defined and see briefly how we can write Python programs that use them

Simplest Regular Expressions

- The simplest regular expressions are those that just contain exactly the characters you are looking for
- So, suppose you want to see if the letters “CSE” appear in a string
- The regular expression in this case would simply be: CSE
- It would match the occurrences of “CSE” in these strings:
 - **CSE**
 - **CSE 101**
 - **I am taking CSE 101**

Digits

- The notation **\d** represents a **single digit**
- So a generic Social Security Number could be expressed as **\d\d\d-\d\d-\d\d\d\d** using regexes
- But, this is a little cumbersome, so we can use **curly braces** instead to denote repetition. For example, **{3}** would mean to repeat something 3 times. Let's rewrite our regex:
- **\d{3}-\d{2}-\d{4}**
- Any SSN would **match** the above pattern
- This is a nifty website that will let you try out regular expressions to see which strings they will match: <http://regex101.com>
- We will use it throughout this lesson

The Dot Metacharacter

- A **period** by itself matches **any character**
- So **.{5}** would match a 5-character sequence of any characters:
 - **12345**
 - **abcde**
 - **a&c(9**
 - **1 3 4** Yes, spaces are matched too!
- If you want to match an actual period, you **escape it**: **\.**
- **\d{3}\.\d{3}\.\d{4}** will match phone numbers written in a form like **010.777.1223**

Matching Specific Characters

- Sometimes the dot metacharacter is too “powerful” because it can match any character
- We can match **specific characters** using regular expressions, by defining them inside **square brackets**
- For example, the pattern **[abc]** will only match a single **a**, **b**, **or c** letter and nothing else

Excluding Specific Characters

- In some cases, we might know that there are **specific characters that we don't want to match**
- For example, we might only want to match phone numbers that are *not* from the area code 880
- To exclude specific characters we use the **square brackets** and the **^ (hat)** character
- For example, the pattern **[^abc]** will match any single character **except for the letters a, b, or c**

Character Ranges

- We can match a character from a list of sequential characters by using the dash to indicate a **character range**
- For example, the pattern **[0-6]** will only match any single digit character from zero to six, and nothing else
- **[^n-p]** will only match any single character except for letters **n** through **p**
- Multiple character ranges can also be used in the same set of brackets
- **[A-Za-z0-9_]** is often used to match characters in English text

Alphanumeric Characters

- The notation **\w** denotes any word character: **a-z, A-Z, 0-9** and the underscore (**_**)
- So the regex **\w{3}\d{4}** would match any string containing three alphanumeric characters followed by any four digits
 - **abc1234**
 - **ABC1234**
 - **1AB7892**
 - **1234567**
- Related is the **\W** notation, which matches any non-alphanumeric character

Repetitions

- We saw that the notation **{m}** means to match the character *m* times
- We can also use the repetition notation with the square bracket notation
- For example, **[wxy]{5}** would match five characters, each of which can be a **w**, **x**, or **y**
- The notation **{m,n}** means we want to match a pattern from **m** to **n** times, inclusive of **m** and **n**
- For example, **.{2,6}** would match from two to six consecutive copies of any character, inclusive
- Another example: **[a-z]{3-7}** would match from three to seven lowercase letters, inclusive

Zero or More Repetitions

- The ***** symbol (the **Kleene star**) means that we want to match **0 or more** repetitions of the character or group of characters that it follows
- The **+** symbol (the **Kleene plus**) matches **1 or more** repetitions of the character or group of characters that it follows
- **\d*** would match any number of digits
- **\d+** would match one or more digits
- **[abc]+** would match one or more of any **a**, **b**, or **c** character in any combination

Optional Characters

- The **?** (question mark) metacharacter denotes **optionality**
- For example, the pattern **ab?c** will match either the strings **abc** or **ac** because the **b** is considered optional
- If you want to match an actual question mark, you *escape it*: **\?**
- The **?** can also be used with groups of characters. More on groups a little later.

Any Whitespace

- Real-world input contains a lot of **whitespace**: spaces, tabs, and newlines
- The whitespace special character **\s** will match any of the specific whitespaces
- Similarly, **\S** will match any non-whitespace character

Starting and Ending Patterns

- We can define a pattern that describes both the start and the end of the line using the special **^** (hat) and **\$** (dollar sign) metacharacters
 - Note that **^** is used in a different context here; when used in square brackets it means to exclude characters
- **^success** would match only a line that begins with the word **success**
- **target\$** would match only a line that ends with the word **target**

Match Groups

- We can use groups of characters and capturing them using the parentheses metacharacters
- For example, imagine that you had a command line tool to list all the PNG image files you have stored on a computer that start with the letters “IMG”
- You could then use a pattern such as **`^(IMG\d+\.png)$`** to capture and extract the full filename

Conditional Matches

- When using groups, you can use the `|` (logical OR, a.k.a. the **pipe**) to denote different possible sets of characters
- We could write the pattern **Buy more (milk | bread | juice)** to match only the strings “Buy more milk”, “Buy more bread”, or “Buy more juice”
- You can use any sequence of characters or metacharacters in a condition
- For example, **([cb]ats* | [dh]ogs?)** would match “cats”, “cat”, “batsssss”, “dog”, “hogs” and many other strings

Regexes in Python

- Python offers some advanced capabilities for working with regexes, but this simple example gives the basics. The **r** character indicates we are using a **raw string**:

```
import re
phone = '123-456-7890'
pattern = r'^\d{3}-\d{3}-\d{4}$'
if re.search(pattern, phone):
    print('The string matches the pattern.')
else:
    print('The string does not match.')
```

Regexes in Python

- Groups are numbered starting with 1 and can be identified by **\1**, **\2**, etc.
- The **sub** function lets us perform replacements in a string
- **sub** takes three arguments:
 - A regex that describes the input string
 - A regex that describes what substitutions to perform
 - The input string itself
- A simple example that replaces **dog** with **fox**:

```
line = 'the dog and cat'  
result = re.sub(r'(.*)(dog)(.*)',  
                r'\1fox\3', line)
```
- **result** will be '**the fox and cat**'

Regexes in Python

- Another example, which reformats a date given in

MM-DD-YY as **YYYY-MM-DD**:

date = '12/25/05'

```
result = re.sub(r'(\d\d)/(\d\d)/(\d\d)',  
                r'20\3-\1-\2', date)
```

- **result** will be '2005-12-25'

Example: Match Filenames

- If you use the Linux or Mac OS terminal frequently, you are often dealing with lists of files
- Most files have a filename component as well as an extension, but in Linux, it is also common to have hidden files that have no filename
 - A hidden file starts with a period
- Let's write a regex to match filenames of only image files (jpg, png, gif) that are not hidden
- Note that the filenames must *end* with one of these three file name extensions
- A regular expression that will match these file names:
- **`\w+\.(png | gif | jpg)`**

Example: Match Filenames

- Test cases (green strings match, red ones don't match):

photo1.jpg

fun_times.png

baby22.gif

1.png

school.jpg

.jpg

..png

Example: Match Credit Cards

- A credit card number is a sequence of 13 to 16 digits, with a few specific digits at the start that identify the card issuer
- Visa card numbers start with a 4 and have 16 digits
- MasterCard numbers start with the numbers 51 through 55 and have 16 digits
- American Express card numbers start with 34 or 37 and have 15 digits
- A single regular expression to match all three kinds of credit cards (Visa, MC, AmEx):
- `(4\d{15}) | (5[1-5]\d{14}) | (3[47]\d{13})`

Example: Match Credit Cards

- Test cases (green strings match, red ones don't match):

4829193238102932

481729382312323

5180129380293232

5484128379172212

5698274398734598

348282019492392

379187239817298

359820938023232

Example: Match ISBN-10

- Suppose a 10-digit ISBN can be written in any of the following sample formats:
 - ISBN 0 93028 923 4
 - ISBN 1-56389-668-0
 - ISBN 1-56389-016-X
- The first digit must be a 0 or 1 and the last symbol can be a digit (0-9) or the capital letter X
- Let's construct a regular expression for matching ISBNs:
- **ISBN [01][\s-]\d{5}[\s-]\d{3}[\s-][0-9X]**

Example: Match ISBN-10

- The regex isn't perfect because it will match strings that mix dashes
- and spaces, such as **ISBN 0-83923 823-0**
- The solution here would be to combine regexes with string processing methods and functions
- More on this issue a little later
- Test cases (green strings match, red ones don't match):

ISBN 0-82918-392-0

ISBN 1 99238 123 X

ISBN 0-83923 823-6

isbn 0-82918-392-0

ISBN-1-33223-233-X

Example: Match Decimal Numbers

- Our goal is to match numbers like these:
 - **3.14529**
 - **-255.34**
 - **128**
 - **1.9e10**
 - **123,340.00**
- As with programming, it's best NOT to try to solve the entire problem in one shot
- How might we solve this problem in an incremental fashion (akin to successive prototyping)?
- What kinds of numbers would be easiest to match first?

Example: Match Decimal Numbers

- Here is how we might try to piece together the regular expression:

\d+

-?\d+

-?\d+(,\d{3})*

-?\d+(,\d{3})*(\.\d+)?

-?\d+(,\d{3})*(\.\d+)?(e-?\d+)?

Example: Match Decimal Numbers

- Test cases (green strings match, red ones don't match):

3.14529

-255.34

128

1.9e10

1.9e-10

-3.1e8

123,340.00

2,55.3

,213.6

.241

Example: Match Phone Numbers

- Suppose we wanted to match phone numbers that could be written in a variety of forms:
 - 1416-555-3456
 - 14165553456
 - 1416 555 3456
 - 1(416) 555-3456
 - 1(416)555-3456
 - 1 416 555 9292
 - 1-416-555-9292
 - 1 (416) 555-3456
- How could we build a single regular expression that will eventually match all these forms (and possibly more)?

-
- Here is how we might try to piece together the regular expression:

`1\d{3}-\d{3}-\d{4}`

`1\d{3}-?\d{3}-?\d{4}`

`1\d{3}[\s-]?\d{3}[\s-]?\d{4}`

`1[\s-]?\d{3}[\s-]?\d{3}[\s-]?\d{4}`

`1[\s-]?(? \d{3})?[\s-]?\d{3}[\s-]?\d{4}`

`1[\s-]?(? \d{3})?[\s-]?\d{3}[\s-]?\d{4}`

Example: Match Phone Numbers

- Test cases (green strings match, red ones don't match):

1416-555-3456

14165553456

1416 555 3456

1(416) 555-3456

1(416)555-3456

1 416 555 9292

1-416-555-9292

1 (416) 555-3456

(416) 555-3456

Example: Match Phone Numbers

- This is not a perfect solution. For example, this regex would accept a phone number like **1(800)-5552468**.
- Sometimes a regular expression that matches only the correct input is very, very complicated.
- In such cases it is usually advisable to use a simpler regex in concert with string processing functions and methods
- In the phone number example, we could ensure that dashes are present or absent in the correct combinations, for instance
- Then we could avoid problems like accepting the phone number at the top of this slide

Questions?
