An Introduction to WEKA 3.9.x

Some of the slides are taken from presentation by Yizhou Sun

1

1

Content

- What is WEKA?
- The Explorer:
  - Preprocess data
  - Classification
  - Clustering
  - Association Rules
  - Attribute Selection
  - Data Visualization
- References and Resources

2

2

What is WEKA?

- **W**aikato **E**nvironment for **K**nowledge **A**nalysis
  - It's a data mining/machine learning tool developed by Department of Computer Science, University of Waikato, New Zealand.
  - Weka is also a bird found only on the islands of New Zealand.
  - https://www.youtube.com/watch?v=1vgA3CN2PH0
  - Weka software is developed in Java.

3

3

Download and Install WEKA

- Website: http://www.cs.waikato.ac.nz/~ml/weka/index.html
- Support multiple platforms (written in java):
  - Windows, Mac OS X and Linux

4

4

Main Features

- 49+ data preprocessing tools
- 76+ classification/regression algorithms
- 8+ clustering algorithms
- 3+ algorithms for finding association rules
- 15+ attribute/subset evaluators + 10+ search algorithms for feature selection

5

5

Main GUI

- Four graphical user interfaces
  - "The Explorer" (exploratory data analysis)
  - "The Experimenter" (experimental environment)
  - "The KnowledgeFlow" (new process model inspired interface)
  - "Workbench" (unified GUI that combines above three)
- One old fashioned Command Line Interface (CLI)

6

6

The package management system

- Weka community keeps adding new algorithms and features.
- These are placed into plugin packages.
- A package management system allows the user to browse and install packages of interest.

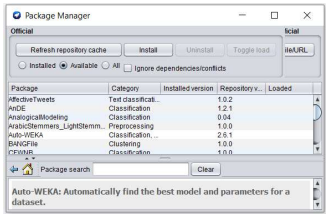

7

7

Content

- What is WEKA?
- The Explorer:
  - Preprocess data
  - Classification
  - Clustering
  - Association Rules
  - Attribute Selection
  - Data Visualization
- References and Resources

8

8

Explorer: pre-processing the data

- Data can be imported from a file in various formats: ARFF, CSV, C4.5, binary
- Data can also be read from a URL or from an SQL database (using JDBC)
- Pre-processing tools in WEKA are called "filters"
- WEKA contains filters for:
  - Discretization, normalization, resampling, attribute selection, transforming and combining attributes, …

9

9

Weka explorer interface

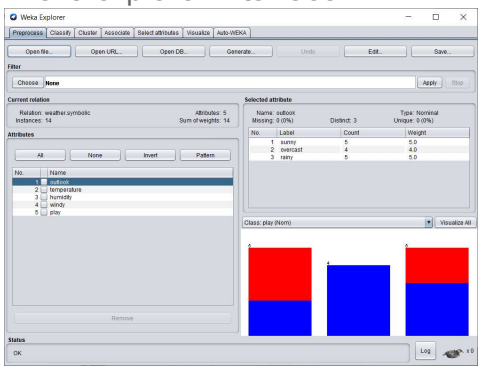

10

10

WEKA: Attribute Relation File Format (ARFF)

CSV format

```
outlook,temperature,humidity,windy,play
sunny,85,85,FALSE,no
sunny,80,90,TRUE,no
overcast,83,86,FALSE,yes
rainy,70,96,FALSE,yes
rainy,68,80,FALSE,yes
rainy,65,70,TRUE,no
overcast,64,65,TRUE,yes
sunny,72,95,FALSE,no
sunny,69,70,FALSE,yes
rainy,75,80,FALSE,yes
sunny,75,70,TRUE,yes
overcast,72,90,TRUE,yes
overcast,81,75,FALSE,yes
rainy,71,91,TRUE,no
```

ARFF format

```
@relation weather

@attribute outlook {sunny, overcast, rainy}
@attribute temperature numeric
@attribute humidity numeric
@attribute windy {TRUE, FALSE}
@attribute play {yes, no}

@data
sunny,85,85,FALSE,no
sunny,80,90,TRUE,no
overcast,83,86,FALSE,yes
rainy,70,96,FALSE,yes
rainy,68,80,FALSE,yes
rainy,65,70,TRUE,no
overcast,64,65,TRUE,yes
sunny,72,95,FALSE,no
sunny,69,70,FALSE,yes
rainy,75,80,FALSE,yes
sunny,75,70,TRUE,yes
overcast,72,90,TRUE,yes
overcast,81,75,FALSE,yes
rainy,71,91,TRUE,no
```

11

11

Load file weather.numeric.arff

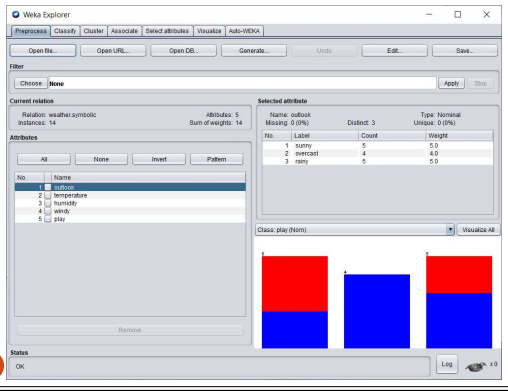

12

12

Building a decision tree – Select J48 (implementation of C4.5 algo)

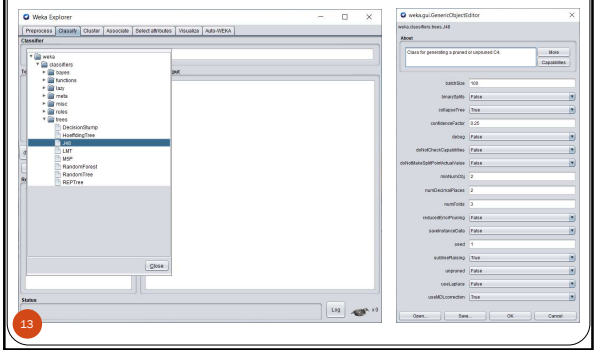

13

13

Using C4.5 classifier

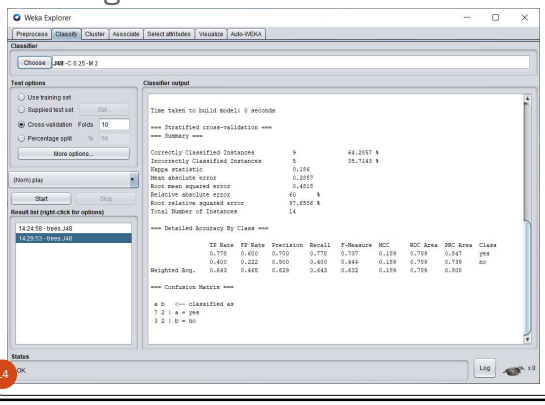

14

14

C4.5 Decision tree visualizer

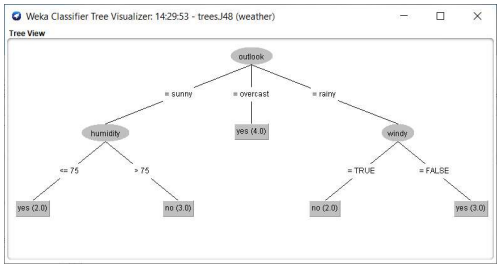

15

15

Using training set for testing

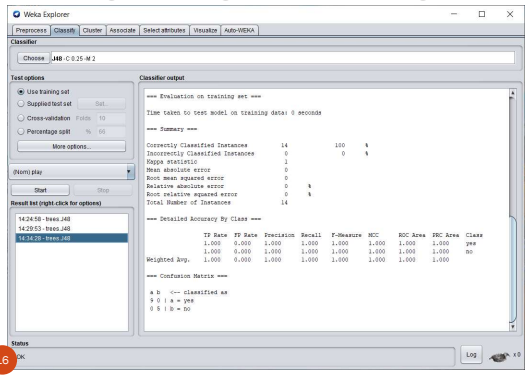

16

16

More options

History line

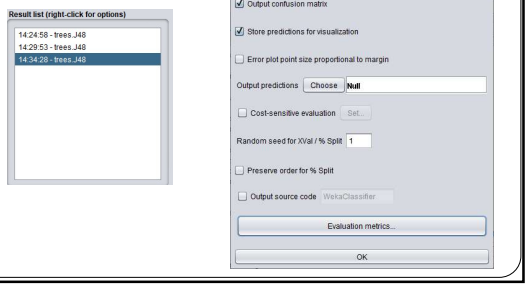

17

17

ROC curve – Visualize threshold curve

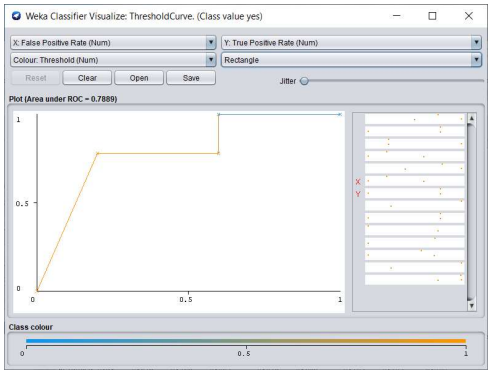

18

18

Working with iris data – iris.arff

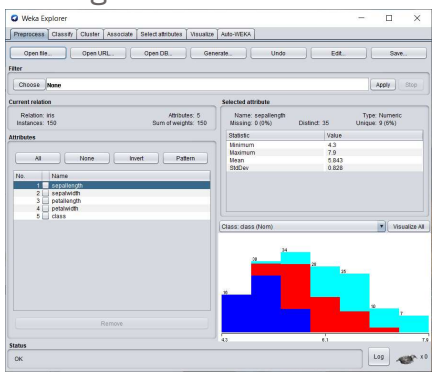

19

19

Use J48 with cross validation

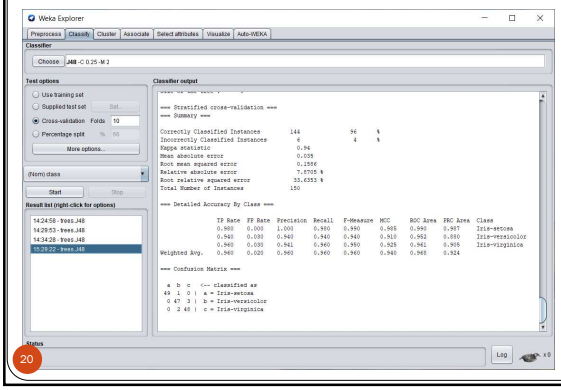

20

20

Iris classifier visualization

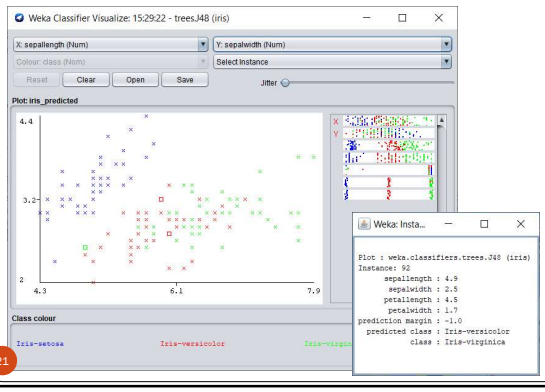

21

21

Iris data – J48 tree

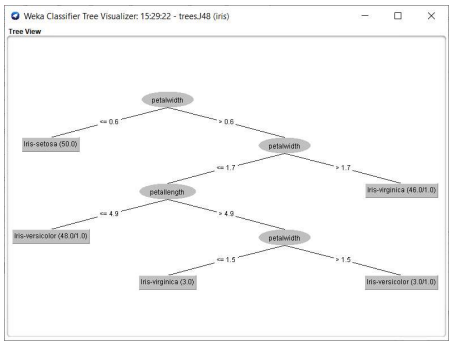

22

22

Open CSV file in Weka

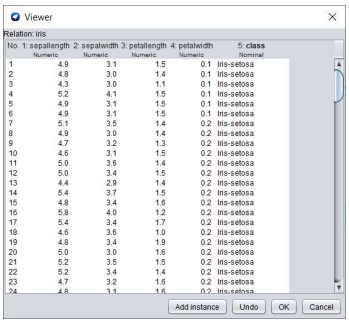

23

23

Filter example – Remove attribute

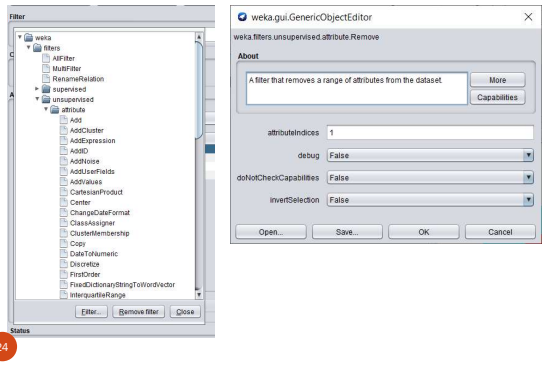

24

24

Iris data with first attribute removed

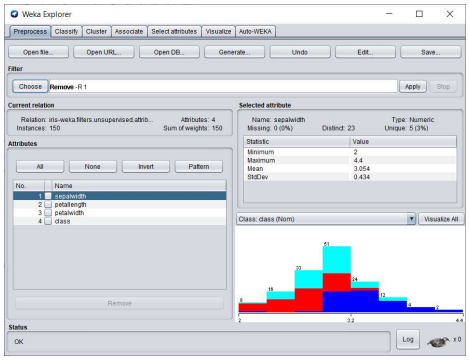

25

25

Load CPU data – cpu.arff

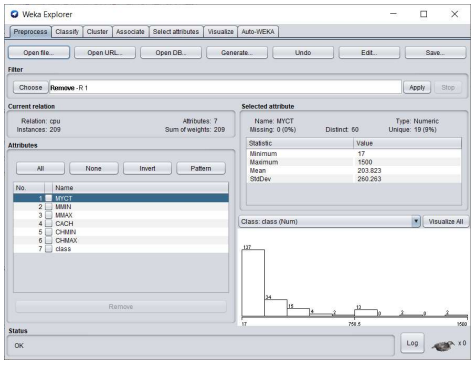

26

26

Choose M5P Quinlan's M5 algorithm Decision trees of regression models

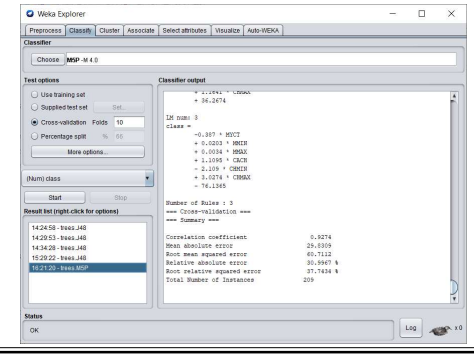

27

27

M5P linear regression models tree

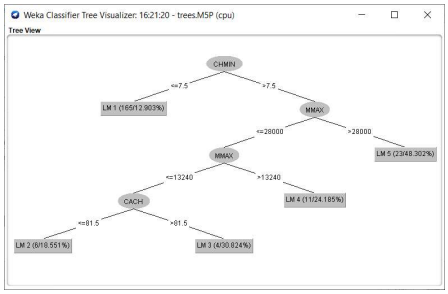

- The first number at each leaf is the number of instances that reach it
- The second is the root mean squared error of the predictions expressed as a percentage of the standard deviation

28

28

Visualization of errors – larger the cross, larger the error – M5P

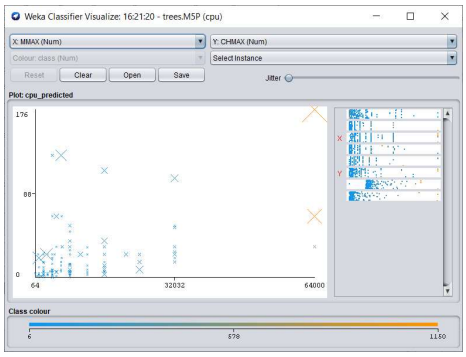

29

29

Choose functions -> Linear regression for CPU data

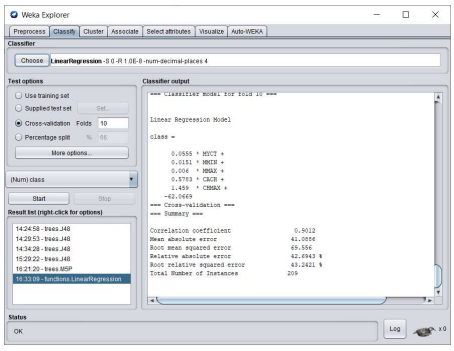

30

30

Visualization of errors –linear regression (M5P is superior)

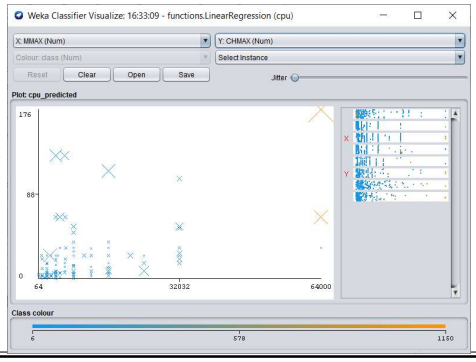

31

31

Naïve Bayes classifier (Iris data)

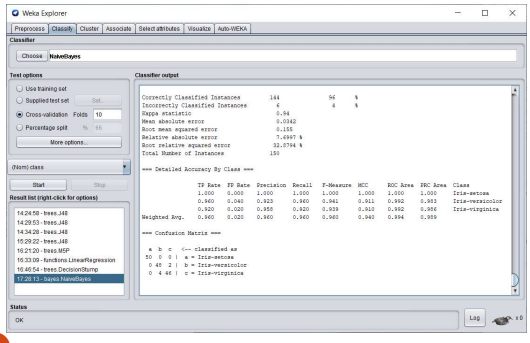

32

32

Naïve Bayes classifier uses normal distribution to model numeric attributes

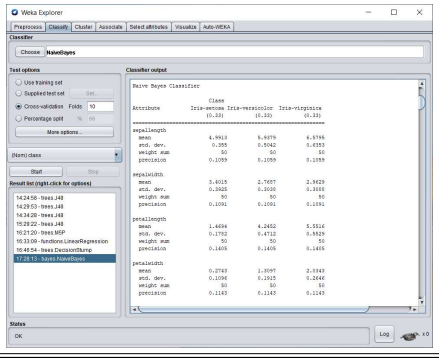

33

33

Naïve Bayes classifier visualization

34

34

Naïve Bayes updatable – process one instance at a time

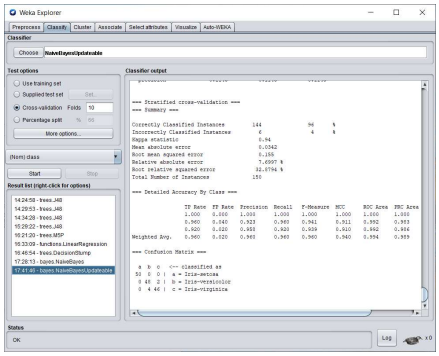

35

35

Naïve Bayes updatable visualization

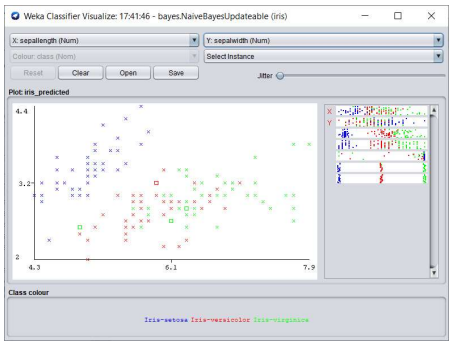

36

36

Explorer: clustering data

- WEKA contains "clusterers" for finding groups of similar instances in a dataset
- Implemented schemes are:
  - *k*-Means, EM, Cobweb, *X*-means, FarthestFirst
- Clusters can be visualized and compared to "true" clusters (if given)
- Evaluation based on loglikelihood if clustering scheme produces a probability distribution

37

37

The K-Means Clustering Method

- Given *k*, the *k-means* algorithm is implemented in four steps:
  - Partition objects into *k* nonempty subsets
  - Compute seed points as the centroids of the clusters of the current partition (the centroid is the center, i.e., *mean point*, of the cluster)
  - Assign each object to the cluster with the nearest seed point
  - Go back to Step 2, stop when no more new assignment

38

38

Clustering – open iris.arff, select Simple Kmeans (numClusters = 3)

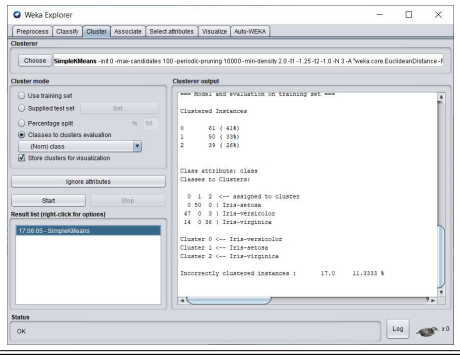

39

39

K-means clustering visualization

40

40

EM (Expectation-Maximization) clustering – out of the box

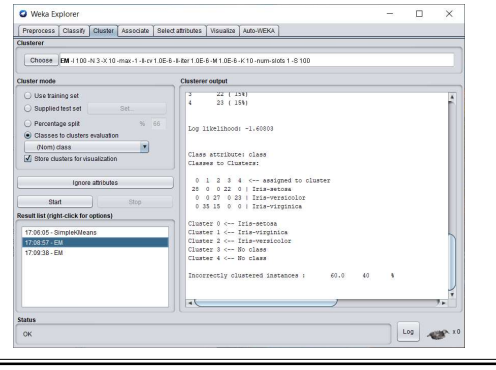

41

41

EM clustering (out of the box) visualization

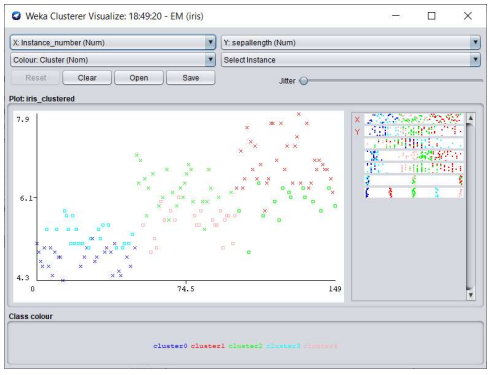

42

42

EM clustering – numClusters = 3

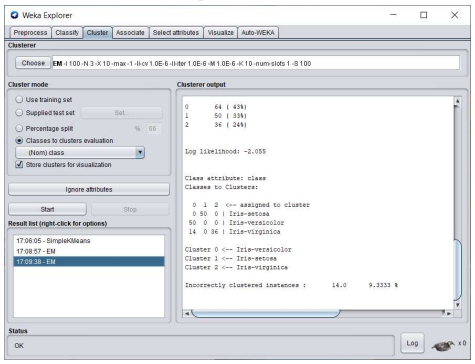

43

43

EM (numClusters = 3) visualization

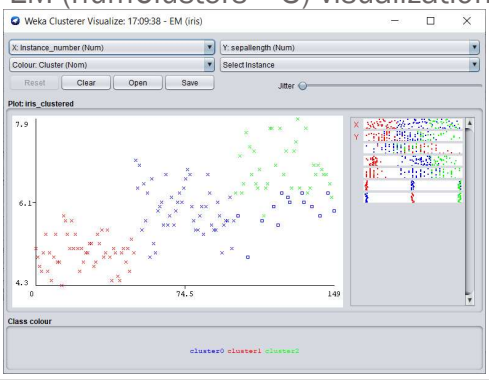

44

44

Hierarchical clustering (numClusters = 3)

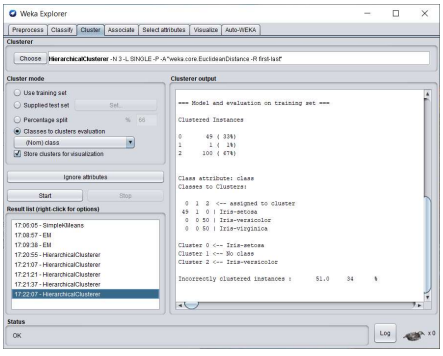

45

45

Hierarchical clustering – tree visualization

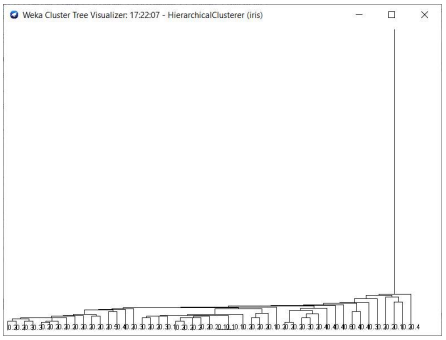

46

46

Hierarchical clustering - visualization

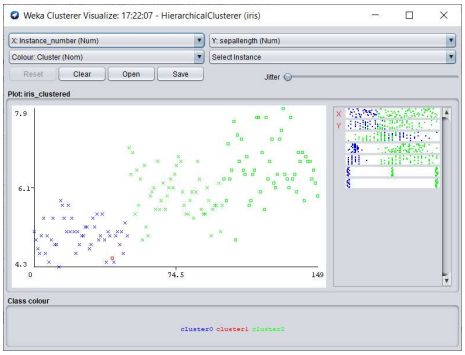

47

47

Explorer: finding associations

- WEKA contains an implementation of the Apriori algorithm for learning association rules
  - Works only with discrete data
- Can identify statistical dependencies between groups of attributes:
  - milk, butter ⇒ bread, eggs (with confidence 0.9 and support 2000)
- Apriori can compute all rules that have a given minimum support and exceed a given confidence

48

48

Basic Concepts: Frequent Patterns

Tid	Items bought
10	Beer, Nuts, Diaper
20	Beer, Coffee, Diaper
30	Beer, Diaper, Eggs
40	Nuts, Eggs, Milk
50	Nuts, Coffee, Diaper, Eggs, Milk

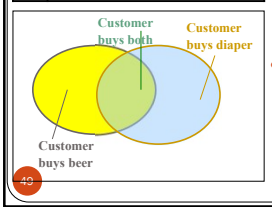

- itemset: A set of one or more items
- k-itemset $X = \{x_1, \ldots, x_k\}$
- *(absolute) support*, or, *support count* of X: Frequency or occurrence of an itemset X
- *(relative) support*, *s*, is the fraction of transactions that contains X (i.e., the probability that a transaction contains X)
- An itemset X is *frequent* if X's support is no less than a *minsup* threshold

49

49

Basic Concepts: Association Rules

Tid	Items bought
10	Beer, Nuts, Diaper
20	Beer, Coffee, Diaper
30	Beer, Diaper, Eggs
40	Nuts, Eggs, Milk
50	Nuts, Coffee, Diaper, Eggs, Milk

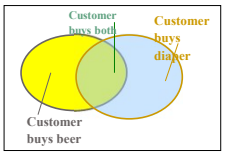

- Find all the rules $X \rightarrow Y$ with minimum support and confidence
  - support, *s*, probability that a transaction contains $X \cup Y$
  - confidence, *c*, conditional probability that a transaction having X also contains *Y*

Let minsup = 50%, minconf = 50%

Freq. Pat.: Beer:3, Nuts:3, Diaper:4, Eggs:3, {Beer, Diaper}:3

- Association rules: (many more!)
  - *Beer* → *Diaper* (60%, 100%)
  - *Diaper* → *Beer* (60%, 75%)

50

50

Associative rules: open supermarket.arff

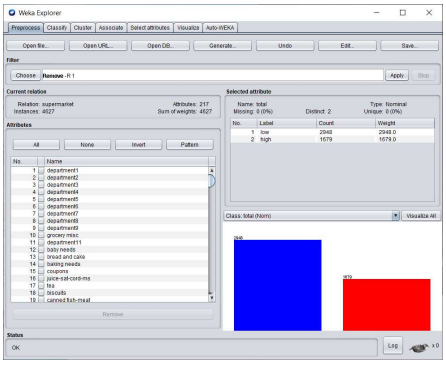

51

51

Associative rules: Apriori algorithm

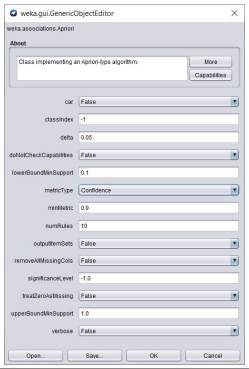

52

52

Associative rules – Apriori output

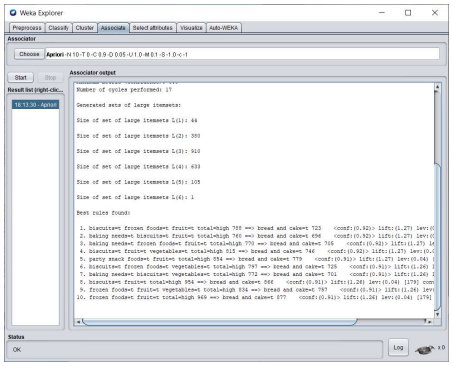

53

53

Weka – Automated feature selection – labor.arff

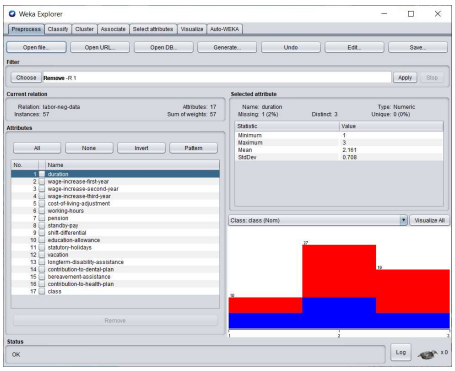

54

54

Weka – Automated feature selection
Select attributes

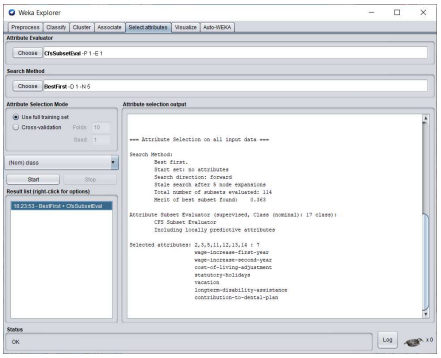

55

55

Attribute selection visualize

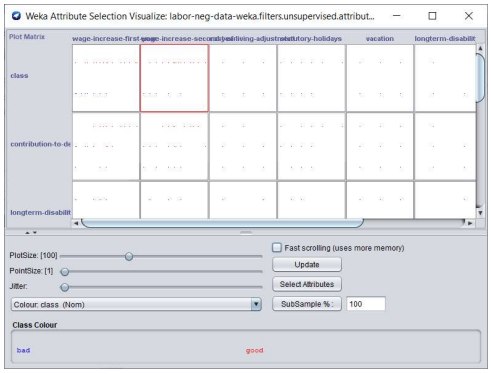

56

56

AutoWeka

- Install AutoWeka plugin from Weka plugin manager
- Automatically searches through the joint space of WEKA's learning algorithms and their respective hyperparameter settings to maximize performance
- Load iris.arff
- Tune AutoWeka parameters
- Hit start

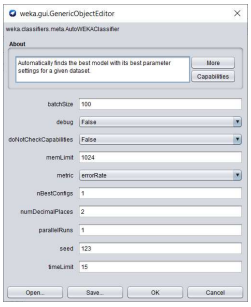

57

57

AutoWeka result

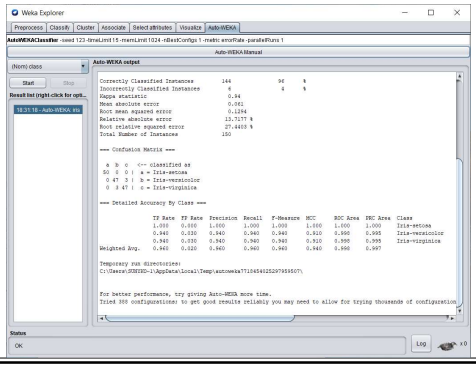

58

58

References and Resources

- References:
  - WEKA website: http://www.cs.waikato.ac.nz/~ml/weka/index.html
  - WEKA Tutorial:
    - Machine Learning with WEKA: A presentation demonstrating all graphical user interfaces (GUI) in Weka.
    - A presentation which explains how to use Weka for exploratory data mining.
  - WEKA Data Mining Book:
    - Ian H. Witten and Eibe Frank, Data Mining: Practical Machine Learning Tools and Techniques (Fourth Edition)
  - WEKA Wiki: http://weka.sourceforge.net/wiki/index.php/Main_Page
  - Auto Weka Software: http://www.cs.ubc.ca/labs/beta/Projects/autoweka/#software
  - Others:
    - Jiawei Han and Micheline Kamber, Data Mining: Concepts and Techniques, 2nd ed.

59

59